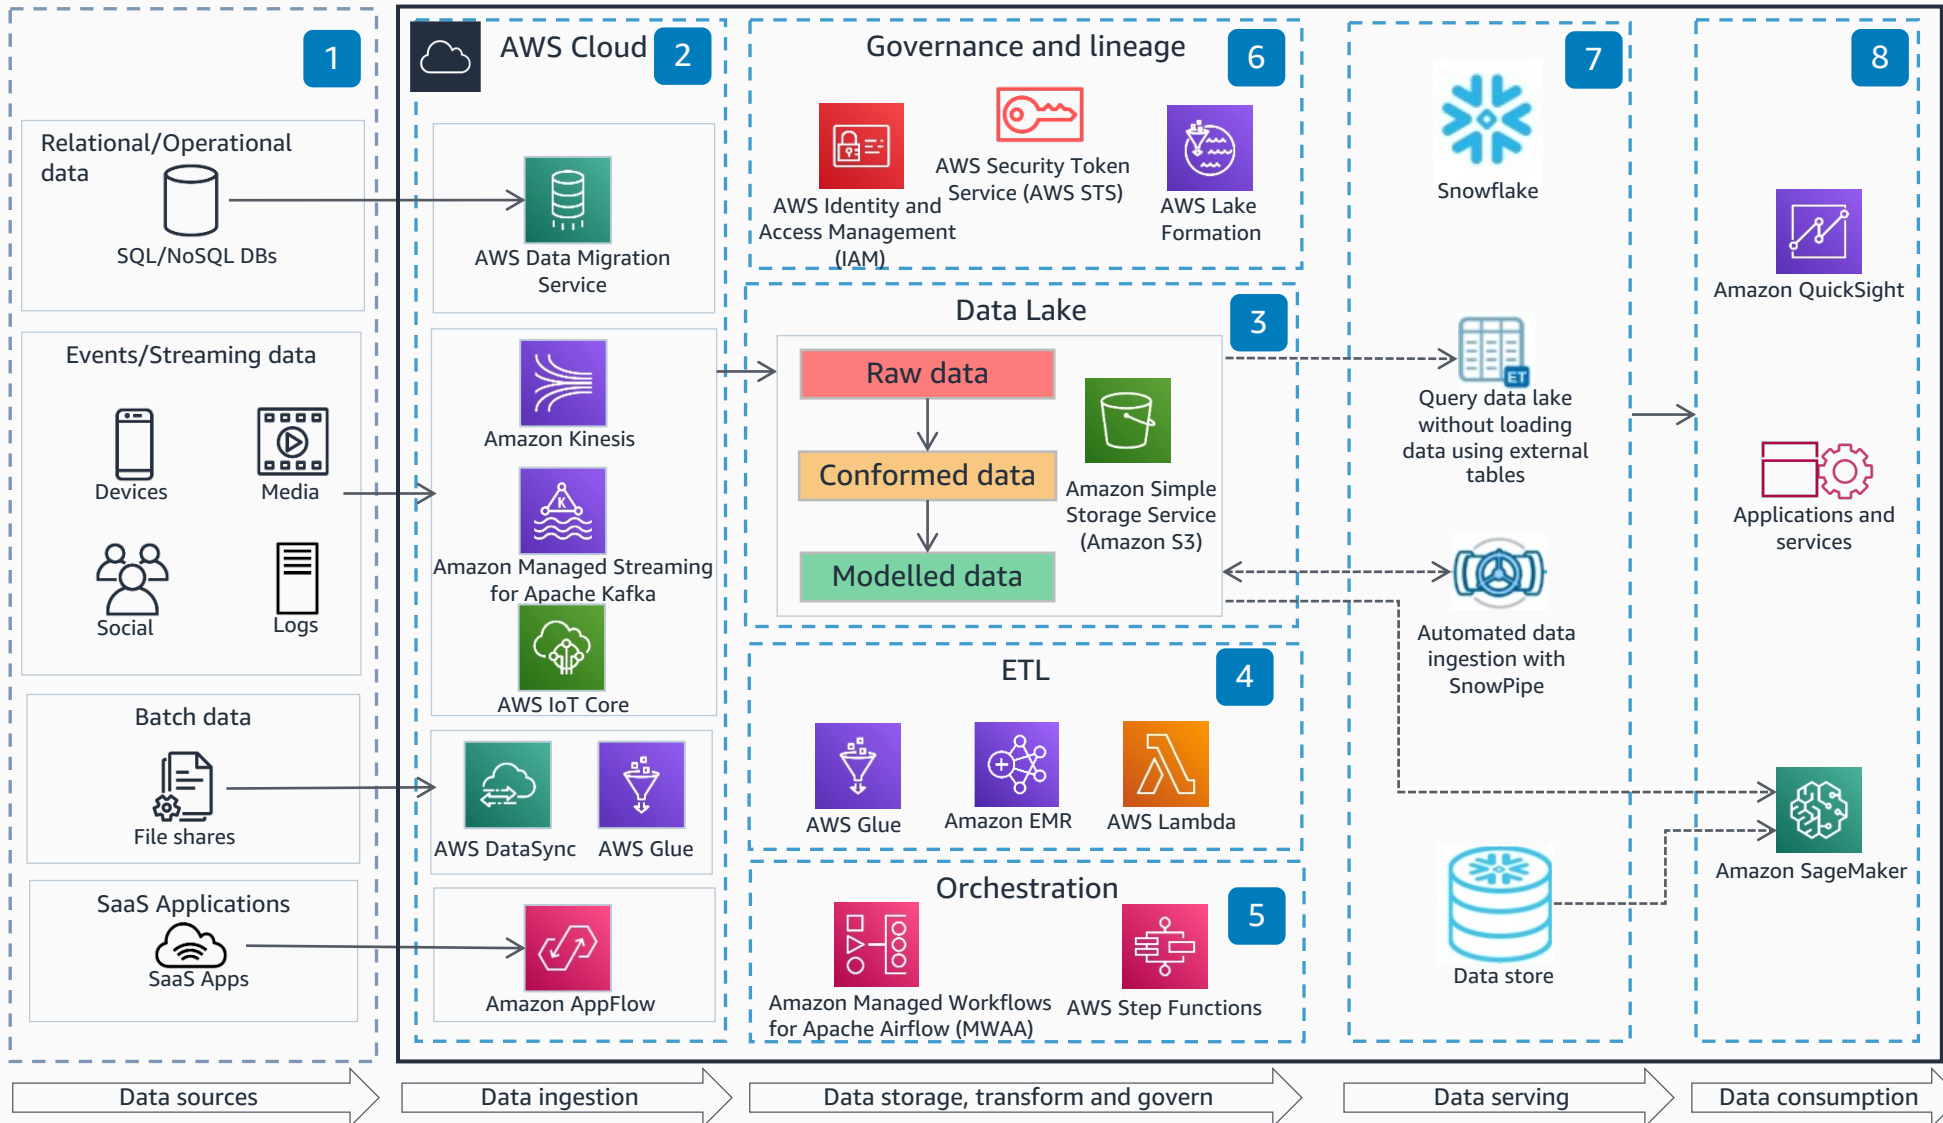


Modern Data Platform using AWS and Snowflake

This architecture enables customers to build end-to-end modern data analytics platforms using AWS and Snowflake.



- 1 Data is collected from multiple data sources across the enterprise, software as a service (SaaS) applications, edge devices, logs, streaming data, and social media networks.
- 2 Based on the type of data source, **AWS Database Migration Service**, **AWS DataSync**, **Amazon Kinesis**, **Amazon Managed Streaming for Apache Kafka**, **AWS IoT Core**, **AWS Glue** and **Amazon AppFlow** are used to ingest the data into the data lake in AWS.
- 3 **Amazon S3** is used for fully managed, highly available and scalable data lake storage.
- 4 **AWS Glue** is used to extract, transform and ingest data across multiple data stores. **Amazon EMR** provides the cloud big data platform for processing vast amounts of data using open source analytics framework. **AWS Lambda** and **Amazon EC2** provide compute capability for data enrichment needs.
- 5 **Amazon Managed Workflows for Apache Airflow (MWAA)** or **AWS Step Functions** is used for orchestrating end-to-end data pipelines.
- 6 **AWS Lake Formation** makes it easy to build, secure and manage your data lake, providing single place to enforce data classification and manage fine-grained access. **AWS IAM** and **AWS STS** provides ability to manage access permissions and temporary credentials.
- 7 Snowflake is used as virtual data warehouse with ability to query **Amazon S3** using external tables, and automated and continuous data ingestion using SnowPipe.
- 8 **Amazon SageMaker** can be used to build, train, and deploy machine learning (ML) models, and add intelligence to your Applications. **Amazon QuickSight** provides ML-powered business intelligence (BI).



Reviewed for technical accuracy March 11, 2022

© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

AWS Reference Architecture